



Stock Market Analysis – Price Movement Prediction

Contents

Abstract

Introduction

Exploratory Data Analysis

Data Cleaning and
transformation
Feature

Creation
Outlier Analysis

Model Performance Analysis

Results and Insights

Conclusion

Abstract

Our project focus is to Predict stock performance within the NIFTY-50 index using machine learning techniques. By studying these price patterns, the project helps uncover trends that make it easier for investors to make smarter financial decisions.

Classification Strategy:

Stocks are categorized into **High**, **Medium**, and **Low** based on price changes.

Stock Data Analysis:

Implemented machine learning models to predict stock performance:

- **Baseline Model:** Serves as a benchmark for comparison.
- **Naïve Bayes**
- **Support Vector Machines (SVM)**
- **Logistic Regression**

Dataset Details

The NIFTY-50 index, represents 50 of India's top companies, serves as a key benchmark for understanding the country's stock market trends.

Dataset contains low frequency stocks data.

The data set contains the stock details from the year 2000 to 2021.

It consists of 470434 rows. 18 columns.

	Date	Symbol	Series	Prev Close	Open	High	Low	Last	Close	VWAP	Volume	Turnover	Trades	Deliverable Volume	%Deliverble
0	2002-02-18	BHARTI	EQ	0.00	51.90	51.90	43.50	44.70	44.35	46.00	10381190.0	4.775431e+13	NaN	6503775.0	0.6265
1	2002-02-19	BHARTI	EQ	44.35	45.15	45.15	41.30	41.35	41.70	42.58	3552660.0	1.512609e+13	NaN	1741184.0	0.4901
2	2002-02-20	BHARTI	EQ	41.70	40.60	42.80	40.25	41.20	41.25	41.56	2512964.0	1.044348e+13	NaN	NaN	NaN
3	2002-02-21	BHARTI	EQ	41.25	42.85	43.40	42.15	42.20	42.40	42.76	1338196.0	5.722681e+12	NaN	485969.0	0.3632
4	2002-02-22	BHARTI	EQ	42.40	42.65	43.60	41.10	43.35	43.30	42.75	811327.0	3.468413e+12	NaN	399133.0	0.4920

Exploratory Data Analysis

Handling Duplicate S

Data Preprocessing and Cleaning

- Drop duplicate rows while keeping the first occurrence

```
Shape before dropping duplicates: (470434, 18)  
Shape after dropping duplicates: (250372, 18)
```

- Renaming stocks that got changed overtime, Replace old stock names with current names in dataframe.

```
['BHARTI' 'BHARTIARTL' 'M&M' 'BRITANNIA' 'TITAN' 'TISCO' 'TATASTEEL' 'LT'  
'HCLTECH' 'BPCL' 'HINDALCO' 'HINDALCO' 'POWERGRID' 'GAIL' 'ITC'  
'BAJAJ-AUTO' 'CIPLA' 'ADANI PORTS' 'ASIANPAINT' 'AXISBANK' 'BAJAJFINSV'  
'BAJFINANCE' 'COALINDIA' 'DRREDDY' 'EICHERMOT' 'GRASIM' 'HDFC' 'HDFCBANK'  
'HEROMOTOCO' 'HINDUNILVR' 'ICICIBANK' 'INDUSINDBK' 'INFRATEL' 'INFY'  
'IOC' 'JSWSTEEL' 'KOTAKBANK' 'MARUTI' 'NESTLEIND' 'NTPC' 'ONGC'  
'RELIANCE' 'SBIN' 'SHREECEM' 'SUNPHARMA' 'TATAMOTORS' 'TCS' 'TECHM'  
'ULTRACEMCO' 'UPL' 'VEDL' 'WIPRO' 'ZEEL' 'UNIPHOS' 'HEROHONDA' 'TELCO'  
'MUNDRAPORT' 'UTIBANK' 'BAJAUTOFIN' 'HINDLEVER' 'INFOSYSTCH' 'JSWSTL'  
'KOTAKMAH' 'SESAGOA' 'SSLT' 'ZEETELE']  
Unique stock values: 66
```

Handling Missing Data

```
Date          0
Symbol        0
Series        0
Prev Close    0
Open          0
High          0
Low           0
Last          0
Close         0
VWAP          0
Volume        0
Turnover      0
Trades        7607
Deliverable Volume  0
%Deliverble   0
dtype: int64
```

Handled missing values using KNN imputation to handle missing values. (3-NN Imputer)

Feature Creation


- **Future_Close:** Calculated by shifting the Close prices forward by a time horizon of 10 trading days.
- $((\text{Future_Close} - \text{Close}) / \text{Close}) \times 100$ to measure the percentage change in stock price over 10 trading days.
- **Price_Movement Classification:**
Categorized Price_Change_% into: High: >4%, Low: < -3%, Medium -3 to 4
- **Class Balancing:** with **stratified sampling 36,172 per class** to avoid bias during model training.
- Applied Label Encoding to transform categorical columns (Symbol and Series) into numerical format.

```
Class Distribution Before Balancing:  
Price_Movement  
Medium      63343  
High        37901  
Low         35597  
Name: count, dtype: int64
```

```
Class Distribution After Balancing:  
Price_Movement  
High         35597  
Low          35597  
Medium       35597  
Name: count, dtype: int64
```

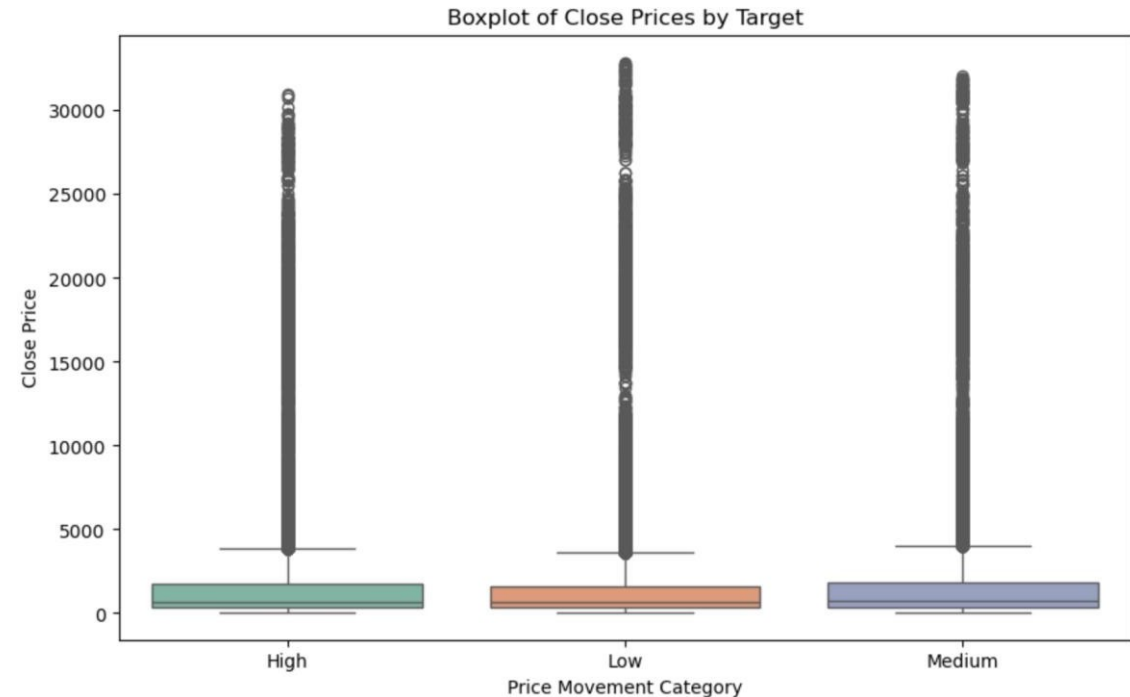

Correlation Analysis Among Input Features

- High correlation among price-related variables, indicating multicollinearity.
- Volume and %Deliverable exhibit weak correlations with price-related variables, indicating they provide unique insights.
- Strong positive correlations exist between Turnover, Trade and Volume, reinforcing their interdependence in transactional metrics.



Distribution For Different Class Labels

- The boxplot displays the distribution of Close prices across three price movement categories: High, Medium, and Low.
- All categories show highly skewed distributions with several outliers, as indicated by the extended whiskers and scattered points.
- Price movements are likely driven by other features, as the spread of Close prices across categories does not reveal significant differentiation.



Outliers Analysis

Number of outliers: 9935

Symbol

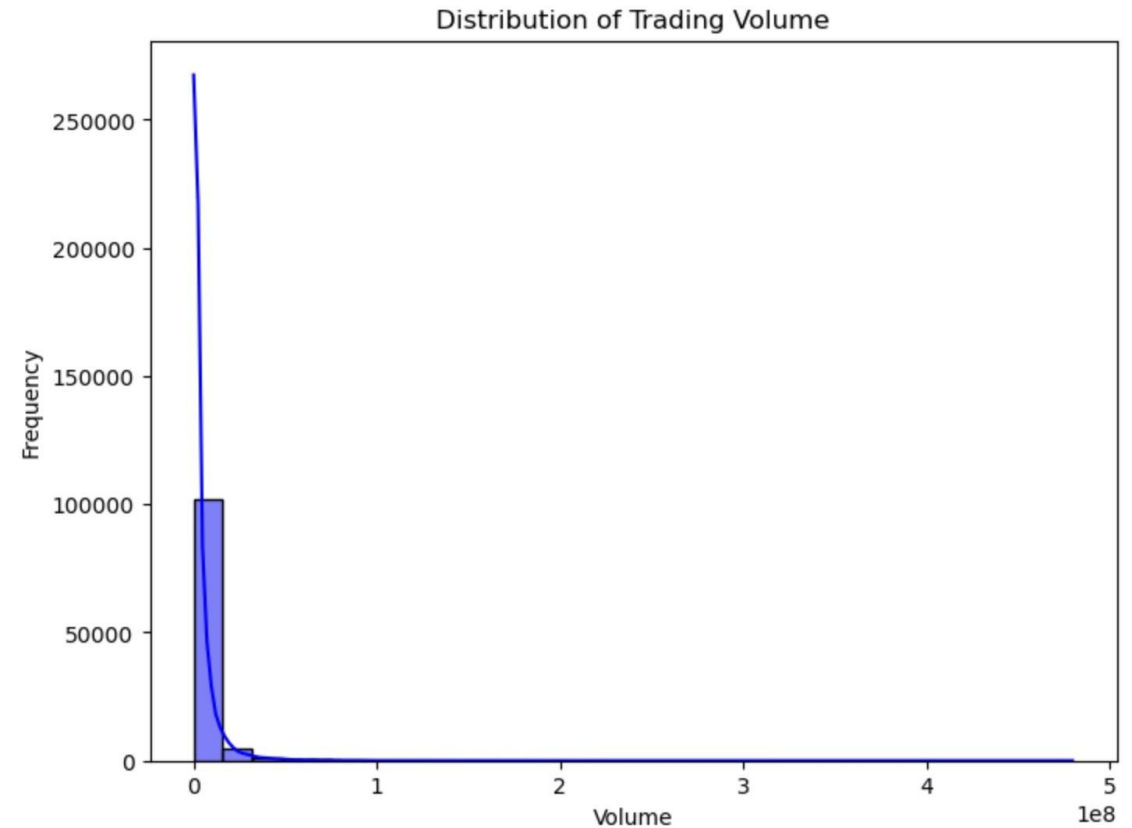
ASIANPAINT	248
BAJAJ-AUTO	42
BAJAJFINSV	931
BAJFINANCE	579
DRREDDY	315
EICHERMOT	1539
GRASIM	184
HEROMOTOCO	108
INFY	66
MARUTI	1293
NESTLEIND	1976
SHREECEM	1709
TITAN	66
ULTRACEMCO	879

dtype: int64

	Symbol	Date	Close	Prev Close	Percent_Change
5	NESTLEIND	2013-01-23	4503.10	4700.75	-4.204648
18	EICHERMOT	2014-02-26	5012.60	4916.15	1.961901
38	BAJAJFINSV	2020-10-23	5831.55	5873.40	-0.712534
59	ULTRACEMCO	2020-11-06	4556.00	4617.90	-1.340436
71	EICHERMOT	2014-11-03	12739.60	12772.75	-0.259537
...
106739	BAJAJFINSV	2020-08-21	6282.40	6289.45	-0.112092
106754	MARUTI	2017-07-26	7565.25	7504.85	0.804813
106760	HEROMOTOCO	2017-07-03	3776.10	3701.35	2.019533
106766	BAJAJFINSV	2021-02-22	10000.60	10250.45	-2.437454
106768	EICHERMOT	2018-04-25	30746.95	31136.75	-1.251897

Distribution For Trading Volume

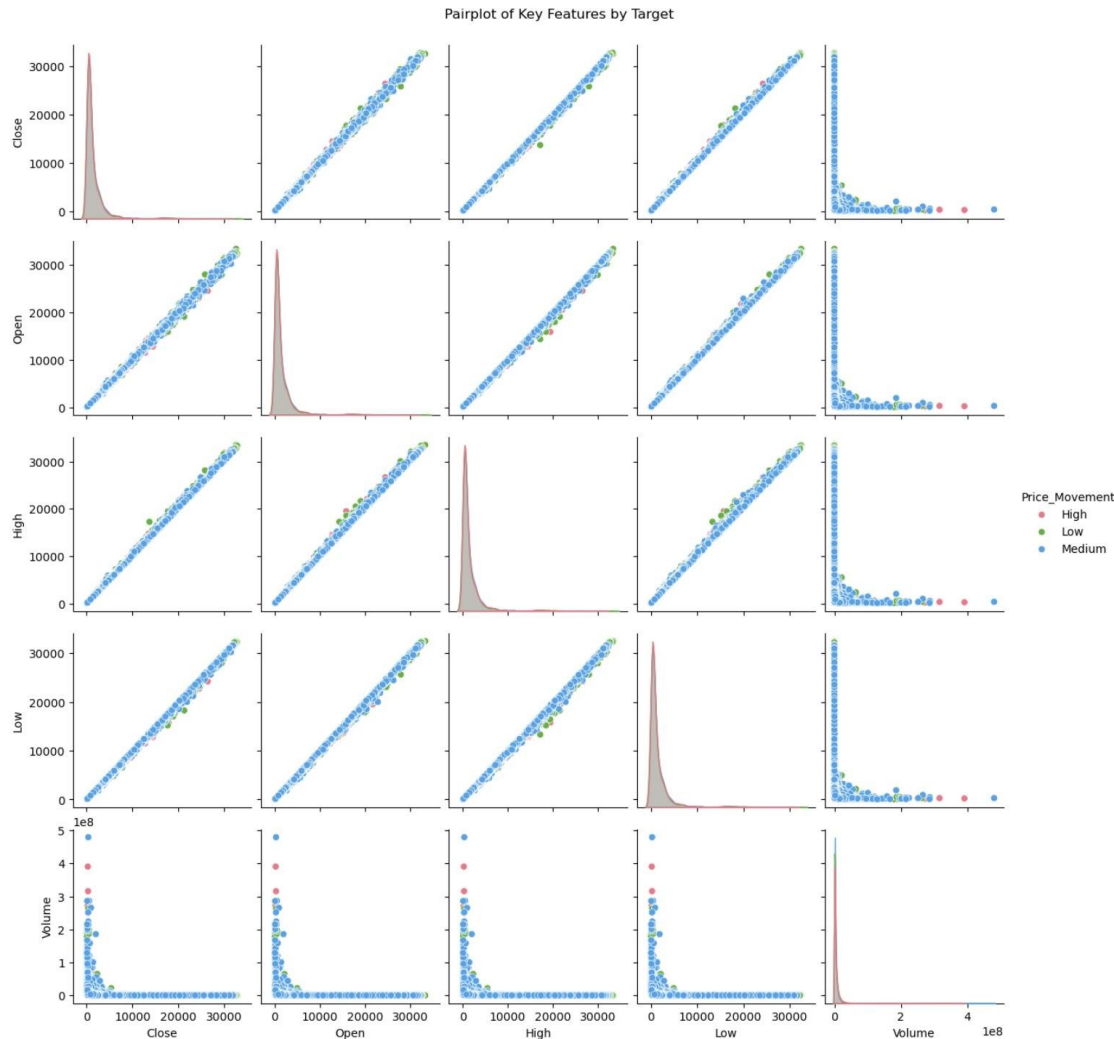
- The Volume distribution is heavily right-skewed, with most data points concentrated in lower ranges.
- Most data points lie in lower ranges, with a few outliers having very high trading volumes.



Pairplots


Pair plot of key features by Target

- Linear relationships among price based features like open, high, low, close.
- Class separability is not evident from pairwise plots.
- Volume shows a distinct pattern, with most data points clustered around low values irrespective of the target class.




Scatterplot of Volume vs. Close Prices by Target

- No clear relationship between Close Price and Volume.
- Significant overlap among target classes.
- Outliers are present in both variables.



PCA Variable Importance And Explained Variance

- The first 5 components explain ~90% of the variance.
- Price-related variables contribute the most to the first component.
- Features like Volume and %Deliverable have lower contributions.



Variable Importance for the First Principal Component:

Low	0.349786
VWAP	0.349780
Close	0.349768
Last	0.349765
Open	0.349753
High	0.349742
Prev Close	0.349682
Future_Close	0.344885
Deliverable Volume	0.093175
Volume	0.091116
Trades	0.076513
Turnover	0.032654
Symbol_encoded	0.024784
%Deliverble	0.011278
Price_Change_%	0.010754

Name: 0, dtype: float64

High and Low Performing Stocks

- High-performing stocks: BajFinance, JSWSteel, and TechM top the list with significant price changes (>5000%).
- Low-performing stocks: LT and Heromotoco show the largest declines in Price_Change_%.
- Recent closing prices for both high and low performers vary significantly, showing diverse market behaviors.

High Performing Stocks (Overall):



	Symbol	dominant_target	Price_Change_%	Close
0	BAJFINANCE	High	7050.231998	4128.40
1	JSWSTEEL	High	6121.189093	886.10
2	INDUSINDBK	High	5556.238104	895.75
3	ADANIPTS	High	5542.158965	319.65
4	TECHM	High	5171.969920	929.60

Low Performing Stocks (Overall):

	Symbol	dominant_target	Price_Change_%	Close
0	LT	Low	-1419.782806	1461.10
1	HEROMOTOCO	Low	-711.972574	2380.15
2	ONGC	Low	-266.038644	263.95
3	GRASIM	Low	-230.690266	2751.10
4	ICICIBANK	Low	1192.680584	288.25

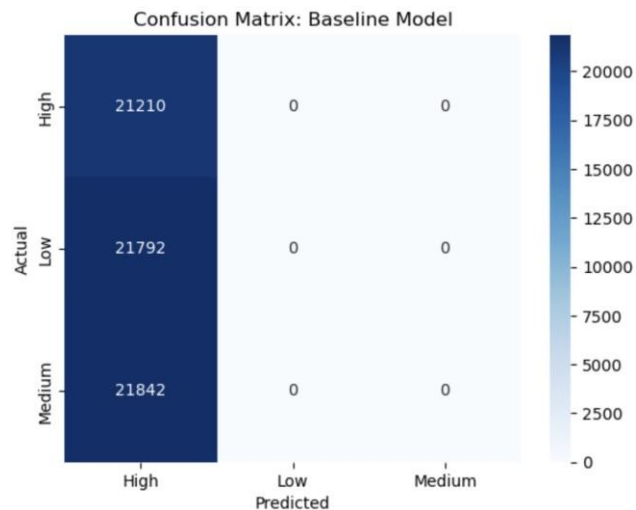
Volatility Analysis and Trends

- Most volatile stocks like Vedl, Gail, and CoalIndia exhibit sharp fluctuations in price trends.
- Least volatile stocks like ITC and TCS show steady and predictable growth.
- Volatility is a key factor in understanding stock performance and risk.



Baseline Model Performance:				
	precision	recall	f1-score	support
0	0.33	1.00	0.49	21210
1	0.00	0.00	0.00	21792
2	0.00	0.00	0.00	21842
accuracy			0.33	64844
macro avg	0.11	0.33	0.16	64844
weighted avg	0.11	0.33	0.16	64844

Accuracy: 0.3270927148232681



Models and Performance Evaluation

Baseline Model


- The baseline model predicts only the majority class (**High**) across all instances, failing to identify the **Low** and **Medium** categories.
- Achieved accuracy of **32.7%**, which is equivalent to predicting the largest class and highlights the need for model improvements.
- **Precision, Recall, and F1-scores** for classes **1 (Low)** and **2 (Medium)** are zero, indicating no correct predictions for these categories.
- All predictions are concentrated in the **High** class, with no instances predicted for other classes, confirming the baseline model's inability to distinguish between categories.

Training Main Model (Naive Bayes)...

Naive Bayes Model Performance:

	precision	recall	f1-score	support
0	0.53	0.41	0.46	21210
1	0.43	0.66	0.52	21792
2	0.53	0.35	0.42	21842
accuracy			0.48	64844
macro avg	0.49	0.48	0.47	64844
weighted avg	0.49	0.48	0.47	64844

Accuracy: 0.47552587749059283



Naïve Bayes Model

- Achieved an accuracy of **47.55%**, a noticeable improvement over the baseline model's performance.
- Performs relatively well for **Low** (class 1) with a recall of **66%**, indicating the model captures a significant portion of these instances.
- Macro and weighted averages for precision, recall, and F1-score are **49%**, reflecting the model's balanced but suboptimal performance.

Tuning Hyperparameters...

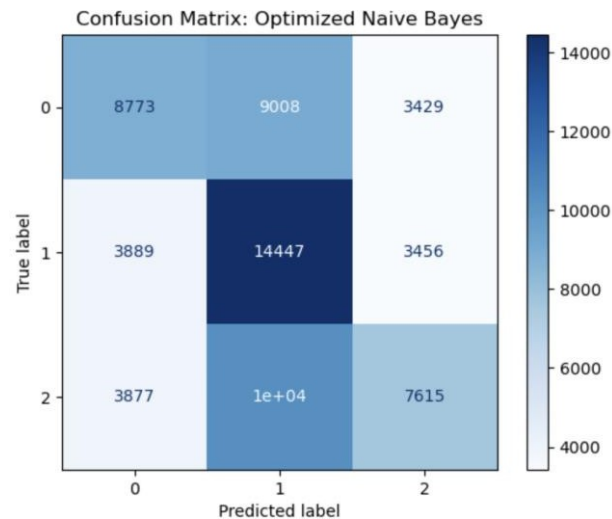
Best Parameters: {'var_smoothing': 1e-10}

Training Optimized Naive Bayes Model...

Optimized Naive Bayes Model Performance:

	precision	recall	f1-score	support
0	0.53	0.41	0.46	21210
1	0.43	0.66	0.52	21792
2	0.53	0.35	0.42	21842
accuracy			0.48	64844
macro avg	0.49	0.48	0.47	64844
weighted avg	0.49	0.48	0.47	64844

Accuracy: 0.47552587749059283



Naïve Bayes Model: Hyper Parameter Tuning

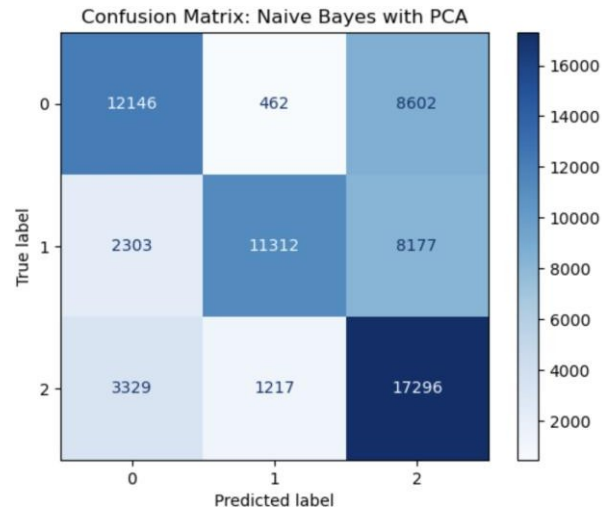
- Best parameter found during tuning: **var_smoothing = 1e-10**
- Achieved an accuracy of **47.55%**, similar to the initial Naive Bayes model.
- Recall for **Class 1 (Low)** remains strong at **66%**, reflecting good detection of this category.
- Misclassifications** are still prominent between **High** and **Low** or **Medium** categories, demonstrating the challenge of clearly separating adjacent classes.

Applying PCA...

Accuracy using PCA-transformed features: 0.6284929985812103

Classification Report for Naive Bayes Model with PCA:

	precision	recall	f1-score	support
0	0.68	0.57	0.62	21210
1	0.87	0.52	0.65	21792
2	0.51	0.79	0.62	21842
accuracy			0.63	64844
macro avg	0.69	0.63	0.63	64844
weighted avg	0.69	0.63	0.63	64844



Naïve Bayes Model: PCA

- Principal Component Analysis (PCA) was applied to reduce feature dimensionality before training the Naive Bayes model.
- Achieved an accuracy of **62.84%**, demonstrating a significant improvement over the previous models.
- Class 1 (Low)** shows the highest precision (**87%**) but lower recall (**52%**), indicating accurate but incomplete predictions.
- Significant misclassifications between **Class 0 (High)** and **Class 2 (Medium)**, indicating areas for further improvement.

Applying PCA...

Tuning Hyperparameters...

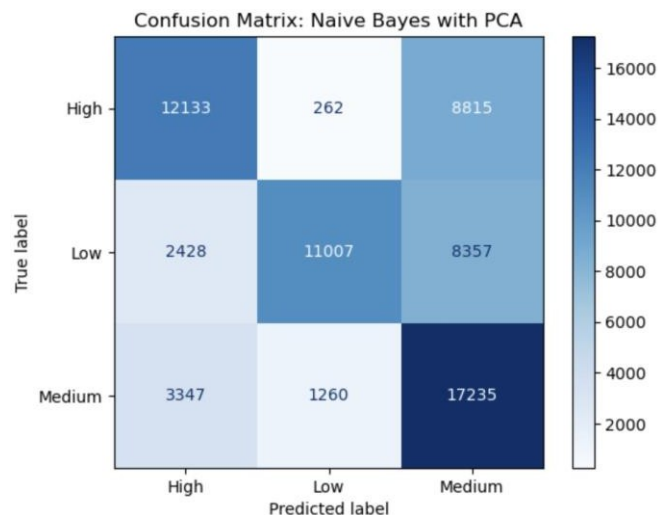
Best Parameters: {'var_smoothing': 2.9763514416313253e-05}

Training Optimized Naive Bayes Model...

Optimized Naive Bayes Model with PCA and Grid Search Performance:


	precision	recall	f1-score	support
0	0.68	0.57	0.62	21210
1	0.88	0.51	0.64	21792
2	0.50	0.79	0.61	21842
accuracy			0.62	64844
macro avg	0.69	0.62	0.62	64844
weighted avg	0.69	0.62	0.62	64844

Accuracy: 0.622648201838258



Naïve Bayes Model: Hyper Parameter Tuning and PCA

- PCA was applied to reduce feature dimensionality, retaining the most critical information for classification and tuning helped refine the model's ability to handle class boundaries effectively.
- Best parameter identified during tuning: **var_smoothing = 2.976e-05**.
- Achieved an accuracy of **62.26%**, slightly lower compared to PCA without tuning (**62.84%**).
- **Class 1 (Low)** retains high precision (**88%**) and relatively balanced recall (**51%**).
- Misclassifications remain evident, particularly between **Class 0 (High)** and **Class 2 (Medium)**.







Testing Data:

	precision	recall	f1-score	support
0	0.95	0.89	0.92	21210
1	0.93	0.89	0.91	21792
2	0.82	0.91	0.86	21842
accuracy			0.90	64844
macro avg	0.90	0.90	0.90	64844
weighted avg	0.90	0.90	0.90	64844

Support Vector Machine Model


- The model achieved an **accuracy of 90%**, demonstrating strong generalization and effective classification.
- **Class 0 (High)** and **Class 1 (Low)** exhibit high precision (**95%** and **93%**) and recall (**89%** each). **Class 2 (Medium)** shows slightly lower precision (**82%**) but excellent recall (**91%**), indicating effective identification despite minor challenges.
- Minor misclassifications are observed between **Class 0 (High)** and **Class 2 (Medium)**.
- AUC scores for all classes range between **0.98–0.99**, validating the model's ability to distinguish among categories effectively.
- Precision-Recall metrics are also robust across all classes.





Testing Data:

	precision	recall	f1-score	support
0	0.99	0.99	0.99	21210
1	0.99	1.00	0.99	21792
2	0.99	0.98	0.98	21842
accuracy			0.99	64844
macro avg	0.99	0.99	0.99	64844
weighted avg	0.99	0.99	0.99	64844



Logistic Regression Model



- The model achieved an outstanding **accuracy of 99%**, demonstrating exceptional generalization capabilities.
- Precision, recall, and F1-scores for all classes are near-perfect at **0.99**, reflecting the model's ability to handle imbalanced data effectively.
- Negligible misclassifications for **Class 0 (High)** and **Class 1 (Low)**, whereas minor confusion observed between **Class 2 (Medium)**.
- AUC scores of **1.00** across all classes validate the model's reliability in distinguishing between categories effectively.
- Near-perfect precision-recall curves indicate a well-calibrated model with consistent performance across varying thresholds.

Logistic Regression Model: Cross Validation

Cross-Validation Results:

```
test_accuracy: Mean = 0.99, Std = 0.00  
test_precision_weighted: Mean = 0.99, Std = 0.00  
test_recall_weighted: Mean = 0.99, Std = 0.00  
test_f1_weighted: Mean = 0.99, Std = 0.00
```

- Cross-validation confirms the model's stability with a **mean accuracy of 99%** and negligible standard deviation across folds.
- Precision, recall, and F1-scores maintain a mean of **0.99** with near-zero variation, showcasing consistent performance.
- These results validate the model's robustness, making it a strong candidate for real-world deployment with minimal risk of overfitting.



Testing Data:				
	precision	recall	f1-score	support
0	0.96	1.00	0.98	21210
1	0.98	1.00	0.99	21792
2	1.00	0.94	0.97	21842
accuracy			0.98	64844
macro avg	0.98	0.98	0.98	64844
weighted avg	0.98	0.98	0.98	64844

Lasso Logistic Regression Model

- The model achieved an **accuracy of 98%**, demonstrating strong predictive performance and generalization capabilities.
- Class 0 (High)** and **Class 1 (Low)** exhibit near-perfect recall (**100%**) and high F1-scores (**0.98** and **0.99**, respectively).
- Negligible misclassifications for **Class 0 (High)** and **Class 1 (Low)**, whereas minor misclassification for **Class 2 (Medium)**, with some overlap into the other categories.
- AUC scores of **1.00** across all classes confirm the model's exceptional ability to distinguish between categories effectively.
- Precision-Recall metrics also remain strong and consistent,


Does Logistic Regression perform well on completely new stock – Britannia ?

	Date	Symbol	Price_Movement	Predicted (Original)
44	2010-11-04	BRITANNIA	Medium	Medium
71	2010-11-05	BRITANNIA	Low	Low
140	2010-11-08	BRITANNIA	Low	Low
187	2010-11-09	BRITANNIA	Low	Low
216	2010-11-10	BRITANNIA	Low	Low
...
139187	2021-04-26	BRITANNIA	Low	Low
139249	2021-04-27	BRITANNIA	Low	Low
139317	2021-04-28	BRITANNIA	Low	Low
139343	2021-04-29	BRITANNIA	Low	Low
139430	2021-04-30	BRITANNIA	Low	Low


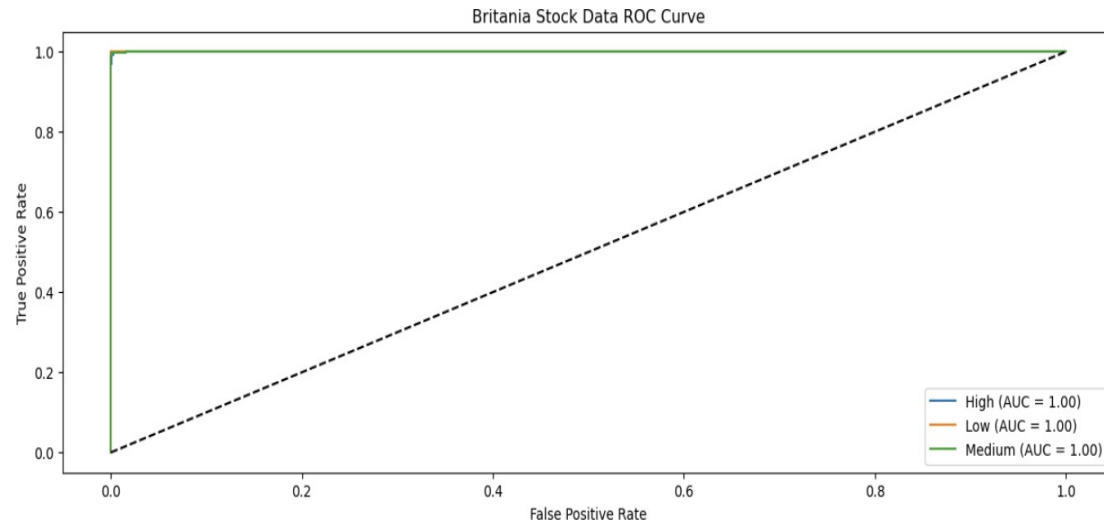
[2598 rows x 4 columns]

Classification Report for Britannia Stock Data:

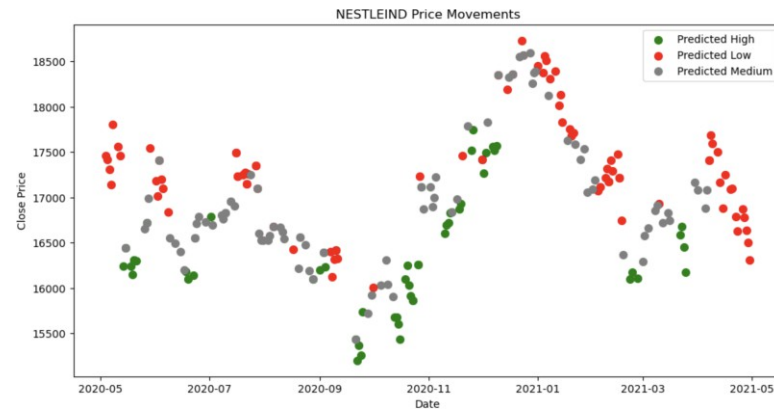

	precision	recall	f1-score	support
0	0.97	1.00	0.99	692
1	0.97	1.00	0.98	511
2	1.00	0.97	0.99	1395
accuracy			0.99	2598
macro avg	0.98	0.99	0.99	2598
weighted avg	0.99	0.99	0.99	2598



Does Logistic Regression perform well on completely new stock – Britannia ?



Visualization of Price Movements Based on Models Classification (For Stocks: Reliance and Nestle)



- The predicted price movements of stocks are visualized over time, categorized as High, Low, or Medium based on the model predictions.
- Scatter plot shows predicted price movements classified as High, Low, and Medium.
- Visualizations validate the model's ability to capture trends in stock price behavior.

Conclusion

The project showcased the importance of iterative model refinement, careful evaluation, and validation in building reliable and interpretable models for complex classification tasks.

Logistic Regression and Lasso Logistic Regression demonstrated exceptional accuracy (98-99%), with balanced precision, recall, and F1-scores across all classes.

Robust ROC and Precision-Recall metrics validated the reliability of the selected models.

Logistic Regression with Lasso is well-suited for deployment in real-world applications due to its scalability and robustness.



Thank

You